Music Recommendation System Using Cosine Similarity

Liang Jimenez

Department of Computer Science, Northeastern University

CS 4100: Artificial Intelligence

Dr. Rose Sloane

December 4, 2024

Introduction

Music consumption has recently become more and more personalized, driven by the availability of tracks. Because of this, the development of recommendation systems that can tailor suggestions to individual user preferences have become increasingly necessary. The challenge in this, however, is in navigating the vast amount of available music to provide relevant and personalized recommendations. This project aims to address this challenge by developing a music recommendation system using a dataset of Spotify tracks, leveraging audio features such as danceability, energy, and loudness to suggest tracks based on user-selected songs.

The foundation of this project is content-based filtering, a method that compares feature vectors of songs to determine similarity. By using cosine similarity, I can quantify how closely related two tracks are based on their audio characteristics and from there, make the best possible recommendation for the user. This approach is effective for high-dimensional data, where traditional methods may fail to capture these nuanced relationships

The work of Zhang (2022) provides insight into the potential of using deep learning techniques for enhancing current music recommendation systems. Zhang's study uses convolutional neural networks (CNNs) to process raw audio data, capturing complex patterns in user preferences and music features. While my project uses a simpler approach because of the pre-extracted features from a Spotify dataset, Zhang's findings highlight the potential for future improvements through deeper learning models.

This project not only aims to create a functional music recommendation system but also explores the machine learning techniques utilized for personalized recommendations. By analyzing methodologies, such as content-based filtering, this study contributes to the broader picture of recommendation systems for music.

Method

This project uses a content-based filtering approach to develop a music recommendation system using a dataset of Spotify tracks. The primary objective is to suggest tracks to users based on a song they select, using audio features such as danceability, energy, and loudness.

Data Collection and Preparation

This project utilizes a publicly available Spotify dataset in CSV format, which includes various audio features for each track. Once the dataset is loaded into a Pandas DataFrame, all unnecessary columns are dropped and all missing values are handled. Additionally, duplicate entries are removed to ensure that each track is uniquely represented. To ensure the data is suitable for analysis, feature vectors for each track are extracted and normalized. This normalization process involves scaling the feature values to a consistent range, which facilitates accurate similarity comparisons.

Implementation of Cosine Similarity

The core algorithm of the recommendation system is based on cosine similarity, a measure used to calculate the cosine of the angle between two non-zero vectors. This method is most effective for high-dimensional data, such as audio feature vectors. These were the steps involved in implementing cosine similarity:

1. Feature Vector Extraction: For each track in the dataset, a feature vector is constructed using selected audio features (danceability, energy, loudness, speechiness, acousticness, intrumentalness, liveness, valence, and tempo).

2. User Selected Song Processing: When a user selects a song of their choice, its feature vector is retrieved and constructed using the same audio features.
3. Similarity Calculation: The cosine similarity between the feature vector of the user-selected song and those of all other tracks in the dataset is commuted using Scikit-learn's '*cosine_similarity*' function.
4. Ranking and Recommendations: Tracks are ranked based on their similarity scores, and the top *N* tracks with highest similarity scores will be shown as recommendations.

Tools and Libraries

The project is coded in Python, using several libraries to execute data manipulation and similarity computation:

- Pandas: Used for data manipulation and preprocessing tasks, such as loading the dataset and normalizing feature vectors
- Scikit-learn: Used to implement cosine similarity calculations

This methodology ensures that the recommendation system can provide personalized music recommendations by effectively comparing audio features across tracks.

Results

The music recommendation system was evaluated based on its ability to suggest tracks that closely match the audio characteristics of a user-selected song. The effective of the system was measured using cosine similarity scores

Cosine Similarity Score Analysis

The system was tested with various user-selected songs across different genres to evaluate its performance. The top 10 recommended tracks were identified based on their cosine similarity scores, which show the degree of similarity between the audio features of the user-selected song and those of other tracks.

Across multiple test cases, the average cosine similarity scores for the top 10 recommendations ranged from 0.85 to 0.95. These high scores indicate that the system effectively identifies tracks with similar audio features to the user-selected song. More specifically, for a user-selected song with high energy and danceability, such as “Uptown Funk” by Mark Ronson featuring Bruno Mars, the system provided recommendations with an average similarity score of 0.98. This proves that the recommended tracks shared similar energetic and danceable features, aligning well with the user’s preferences

These results demonstrate that using cosine similarity to compare feature vectors is an effective method for generating personalized music recommendations. The high similarity scores across various test cases confirm that the system can consistently provide relevant track suggestions based on user input.

Conclusion

This music recommendation system has shown the potential of content-based filtering in providing personalized music recommendations. By using audio features such as danceability, energy, and loudness, the system effectively identifies tracks that align with the characteristics of a user-selected song. The evaluation results, with cosine similarly scores consistently ranging from 0.85-0.95 for the top recommendations, confirm the system’s ability to provide relevant and coherent recommendations

Strengths and Limitations

One of the primary strengths of this approach is the simplicity and efficiency. The use of pre-extracted audio features allows for rapid computation and a relatively straightforward implementation, making it possible for real-time applications. This approach is most effective for high-dimensional data, allowing for straightforward implementation and scalbiity. Additionally, the use of cosine similarity is well-suited for content-based filtering, as it effectively measures the closeness between feature vectors, ensuring relevant track suggestions based on user-selected songs. This method, however, also presents limitations specifically in capturing more complex patterns in user preferences and music features. Unlike advanced models like convolution neural networks, which can process raw audio data to find relationships, the current system relies on predefined features that may not fully encapsulate all of the musical content. As Zhang's study highlights, CNNs can significantly enhance recommendation accuracy by learning from historical user behavior and extracting richer repreentations from audio data.

Future Improvements

To enhance the system's performance and address its limitations, future improvements that explore deep learning techniques similar to those discussed in Zhang's study on CNN-based music recommendation systems could be implemented. By incorporating models capable of processing raw audio data, it may be possible to achieve a deeper understanding of user preferences and improve recommendation accuracy. Additionally, the current system's dataset is limited to the tracks available at the time of data collection, potentially missing newer releases or less popular genres that coud enrich the user experience. Incorporating Spotify's Web API could significantly expand the database available for recommendations, allowing access to up-to-date track features and metadata. This integration would enable the system to provide more diverse and comprehensive music suggestions.

In conclusion, while the current project successfully demonstrates the viability of content-based filtering for music recommendations, it also shows opportuiens for growth through more complex machine learning approaches and expanded data sources. These advancements could singiificanly enhance personalization and accuracy in music recommendation systems.

References

Zhang, Y. (2022). *Music Recommendation System and Recommendation Model Based on Convolutional Neural Network.* Wiley Online Library.